Where does the error come from?

Bias and Variance of Estimator

- Estimate the mean of a variable x
  - assume the mean of x is μ
  - assume the variance of x is σ^2
- Estimator of mean μ
  - Sample N points:$\{x^1, x^2, \ldots, x^N\}$

$$m = \frac{1}{N}\sum_n x^n \neq \mu$$

- Estimator of variance σ^2
  - Sample N points:$\{x^1, x^2, \ldots, x^N\}$

$$m = \frac{1}{N}\sum_n x^n$$

$$s = \frac{1}{N}\sum_n (x^n - m)^2$$

$$\mathrm{Var}[m] = \frac{\sigma^2}{N}$$

Biased estimator

$$E[s] = \frac{N-1}{N}\sigma^2 \neq \sigma^2$$

Variance

$y = b + w \cdot x_{cp}$

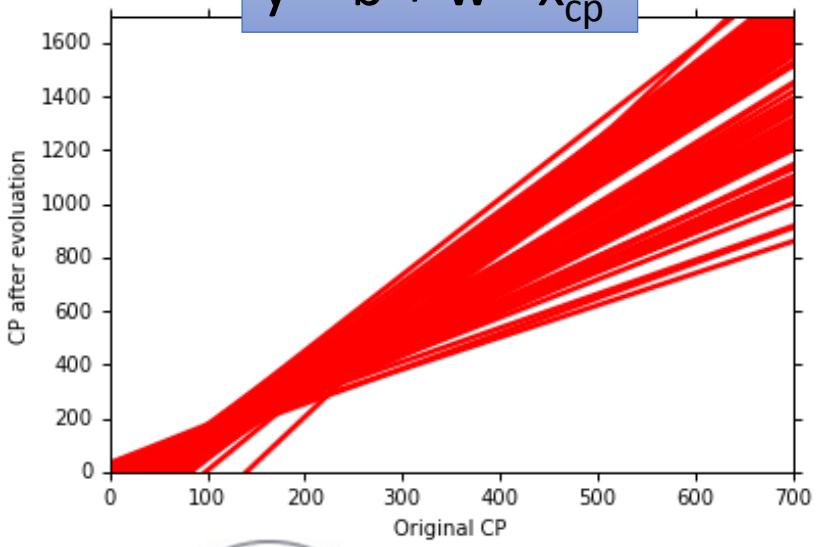

Small Variance

$y = b + w_1 \cdot x_{cp} + w_2 \cdot (x_{cp})^2 + w_3 \cdot (x_{cp})^3 + w_4 \cdot (x_{cp})^4 + w_5 \cdot (x_{cp})^5$

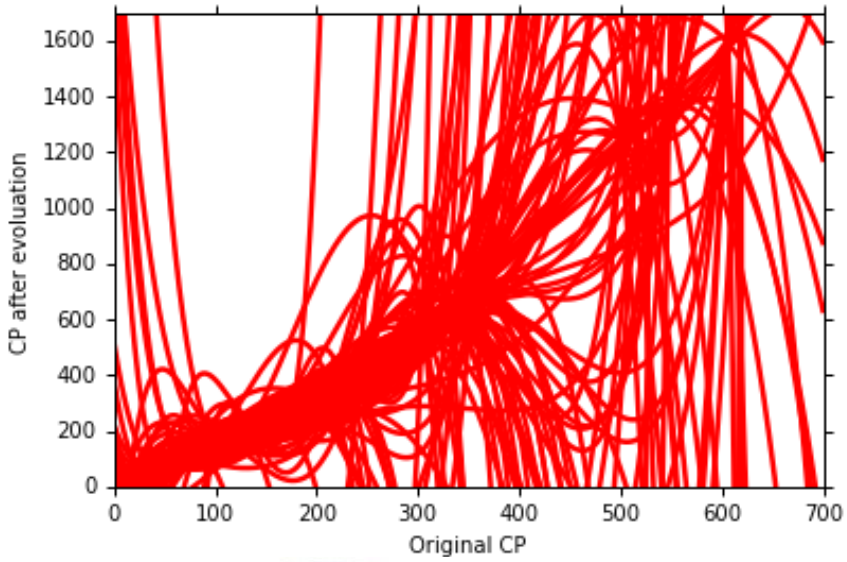

Large Variance

Simpler model is less influenced by the sampled data

Bias

$$E[f^*] = \bar{f}$$

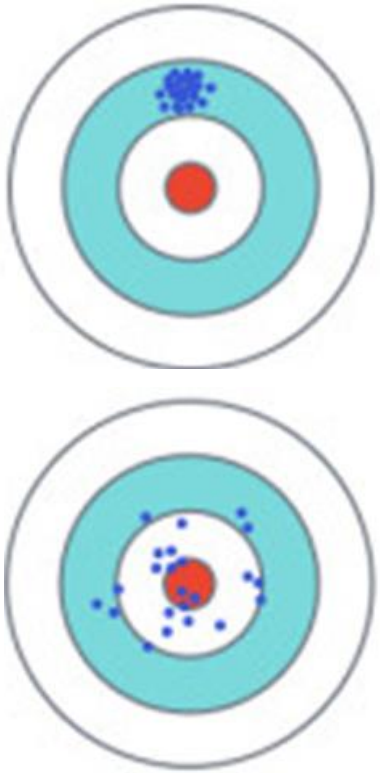

Large Bias

Small Bias

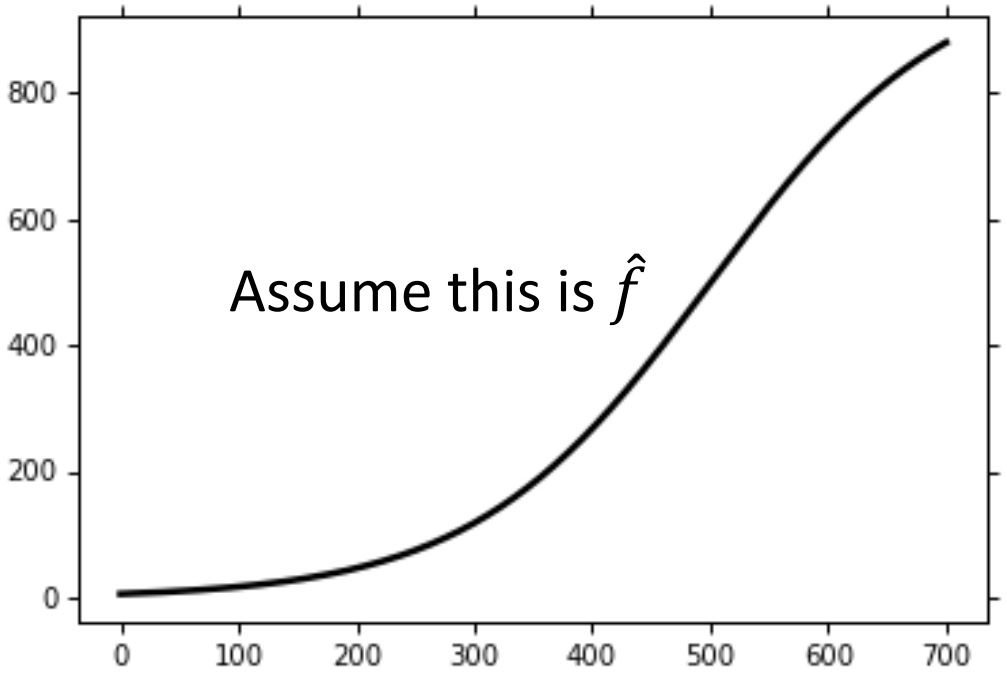

What to do with large bias?

- Diagnosis:
  - If your model cannot even fit the training examples, then you have large bias Underfitting
  - If you can fit the training data, but large error on testing data, then you probably have large variance Overfitting

- For bias, redesign your model:
  - Add more features as input
  - A more complex model

What to do with large variance?

- More data

Very effective, but not always practical

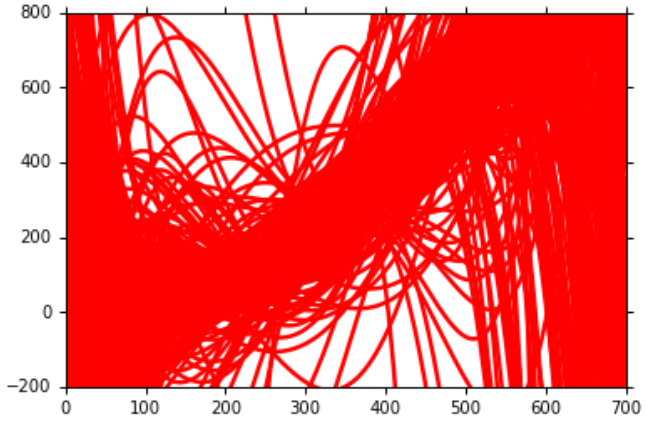

10 examples

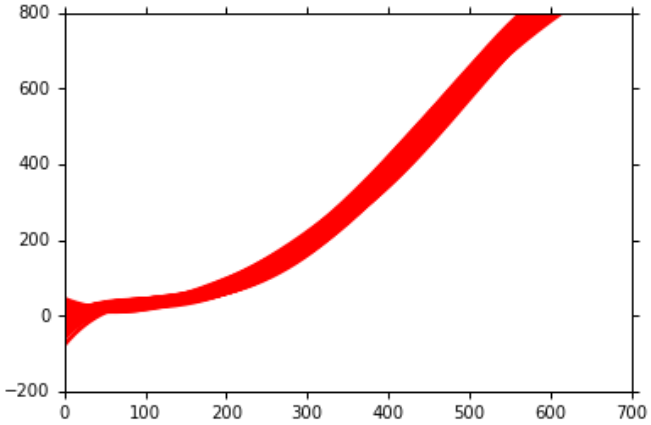

100 examples

- Regularization ➡ May increase bias

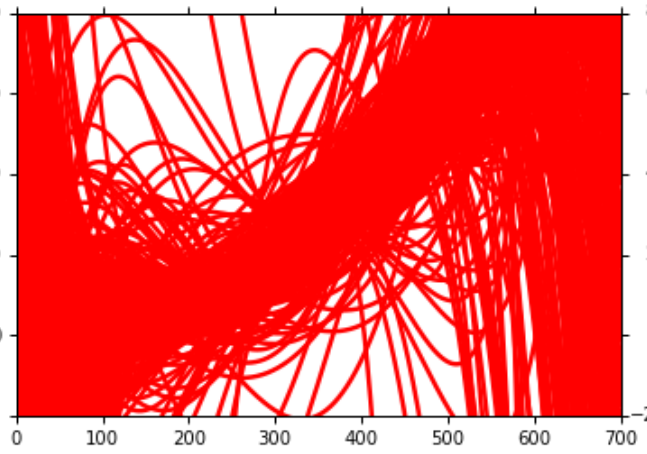

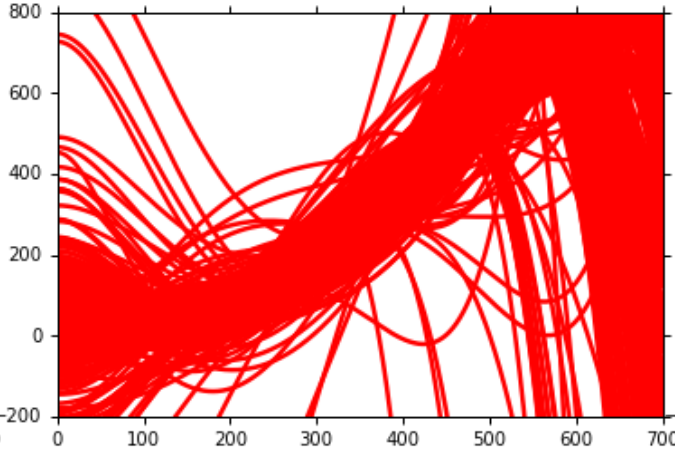

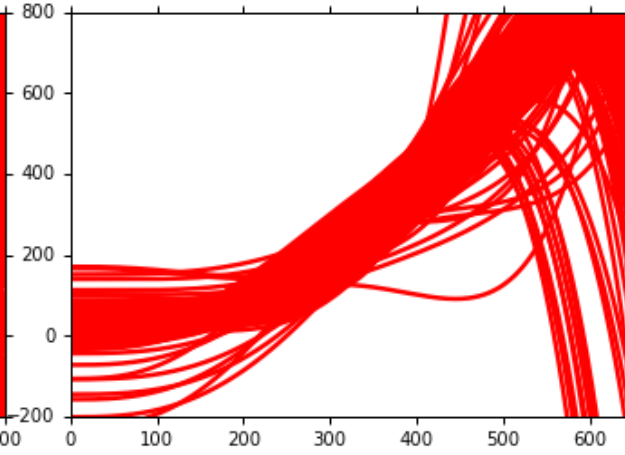

Model Selection

- There is usually a trade-off between bias and variance.
- Select a model that balances two kinds of error to minimize total error

Cross Validation

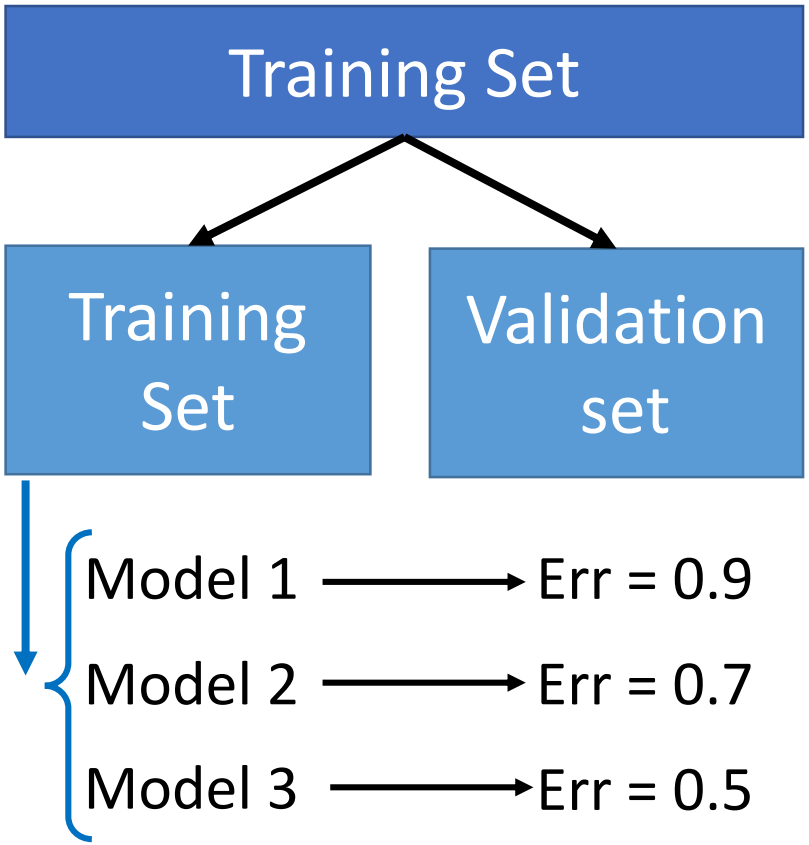

Using the results of public testing data to tune your model
You are making public set better than private set.

N-fold Cross Validation