

“Hello world”
of deep learning

Keras


Step 1:
define a set
of function



Step 2:
goodness of
function



Step 3: pick
the best
function



```
model = Sequential()
```

```
model.add( Dense( input dim=28*28,  
                  output dim=500 ) )  
model.add( Activation('sigmoid') )
```

softplus, softsign, relu, tanh,
hard_sigmoid, linear

```
model.add( Dense( output dim=500 ) )  
model.add( Activation('sigmoid') )
```

```
model.add( Dense(output_dim=10 ) )  
model.add( Activation('softmax') )
```

Keras



Step 3.1: Configuration

```
model.compile(loss='categorical_crossentropy',  
              optimizer='adam',  
              metrics=['accuracy'])
```

SGD, RMSprop, Adagrad, Adadelata, Adam, Adamax, Nadam

Step 3.2: Find the optimal network parameters

```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

Training data
(Images)


Labels
(digits)

In the following slides


We do not really minimize total loss!

Mini-batch

Mini-batch



Mini-batch



➤ Randomly initialize network parameters

➤ Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once

➤ Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once

:

➤ Until all mini-batches have been picked

one epoch


Repeat the above process

Mini-batch

Batch size influences both ***speed*** and ***performance***. You have to tune it.

```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

Mini-batch



100 examples in a mini-batch


Batch size = 1 ➡

Stochastic gradient descent

- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once
- Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once
⋮
- Until all mini-batches have been picked

Speed • Smaller batch size means more updates in one epoch

Keras



Save and load models

<http://keras.io/getting-started/faq/#how-can-i-save-a-keras-model>

How to use the neural network (testing):

case 1:

```
score = model.evaluate(x_test,y_test)
print('Total loss on Testing Set:', score[0])
print('Accuracy of Testing Set:', score[1])
```

case 2:

```
result = model.predict(x_test)
```

Keras

- Using GPU to speed training
 - Way 1
 - `THEANO_FLAGS=device=gpu0 python YourCode.py`
 - Way 2 (in your code)
 - `import os`
 - `os.environ["THEANO_FLAGS"] = "device=gpu0"`