Why Deep Learning?

Modularization

- Deep → Modularization

Don’t put everything in your main function.

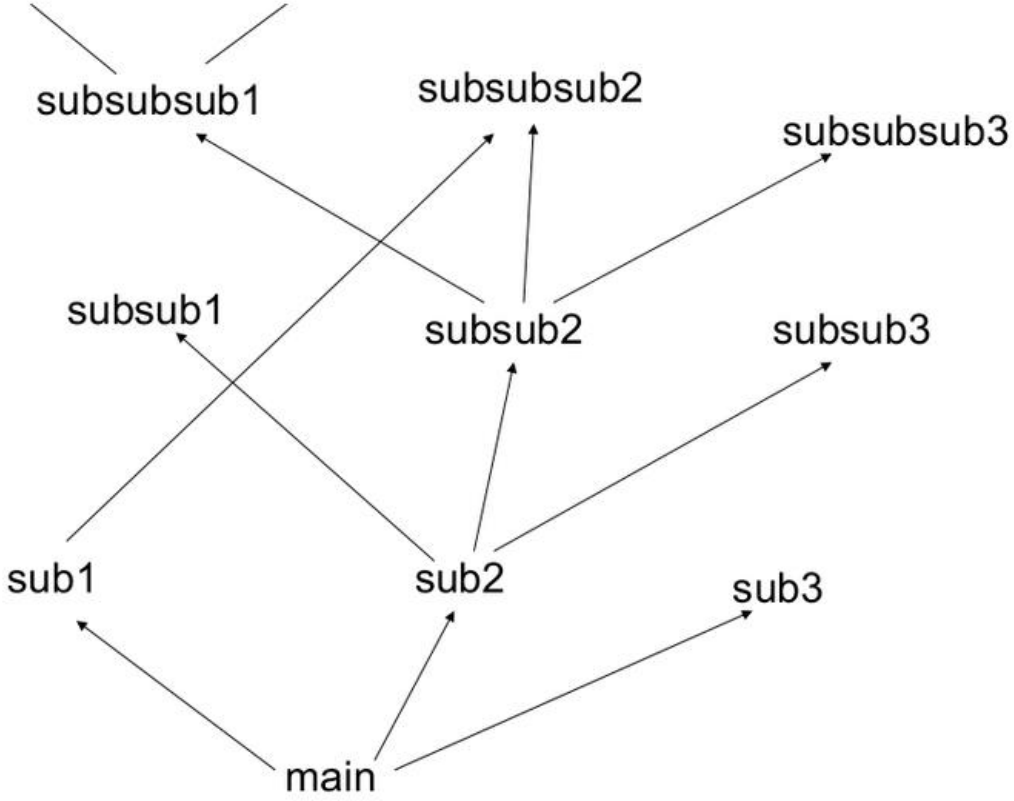

http://rinuboney.github.io/2015/10/18/theoretical-motivations-deep-learning.html

Modularization

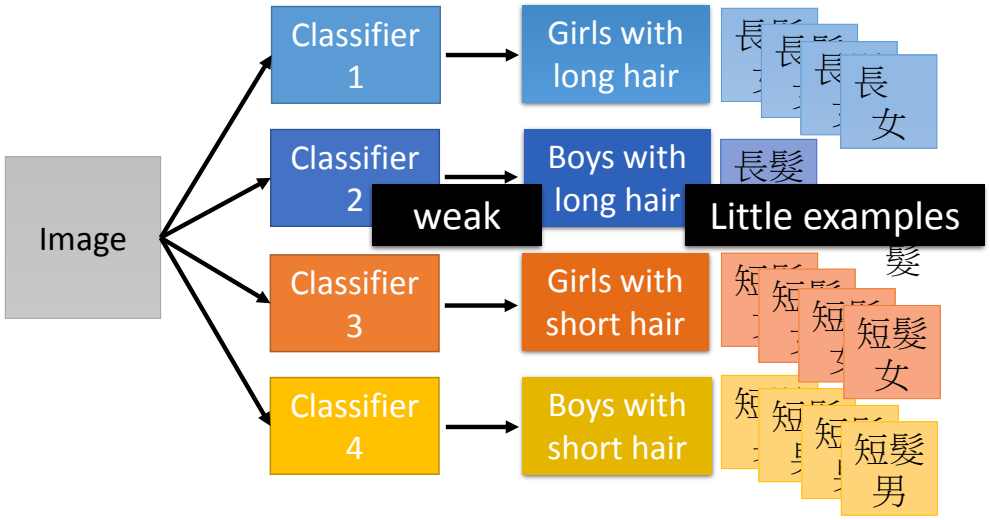

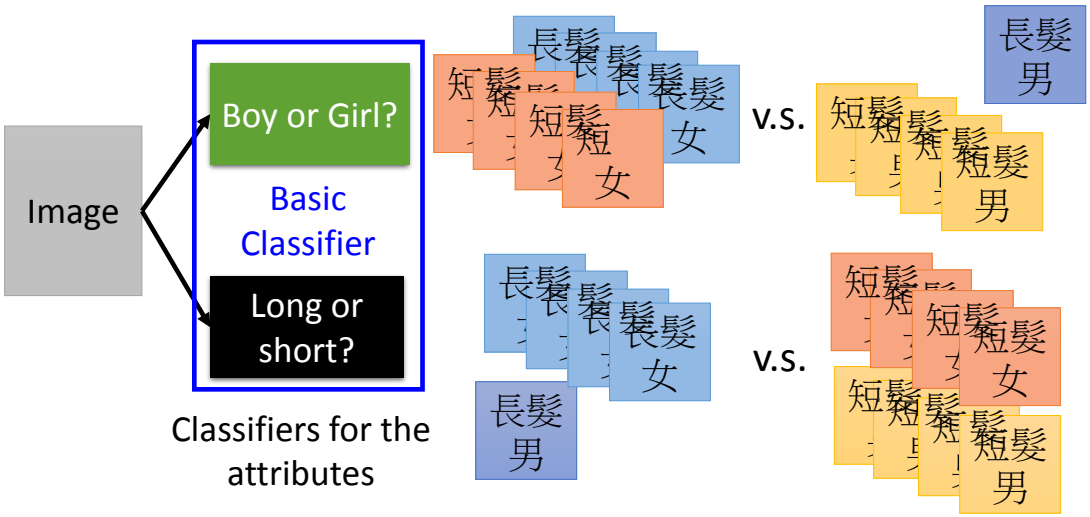

- Deep→Modularization
- Each basic classifier can have sufficient training examples.

Modularization

- Deep → Modularization

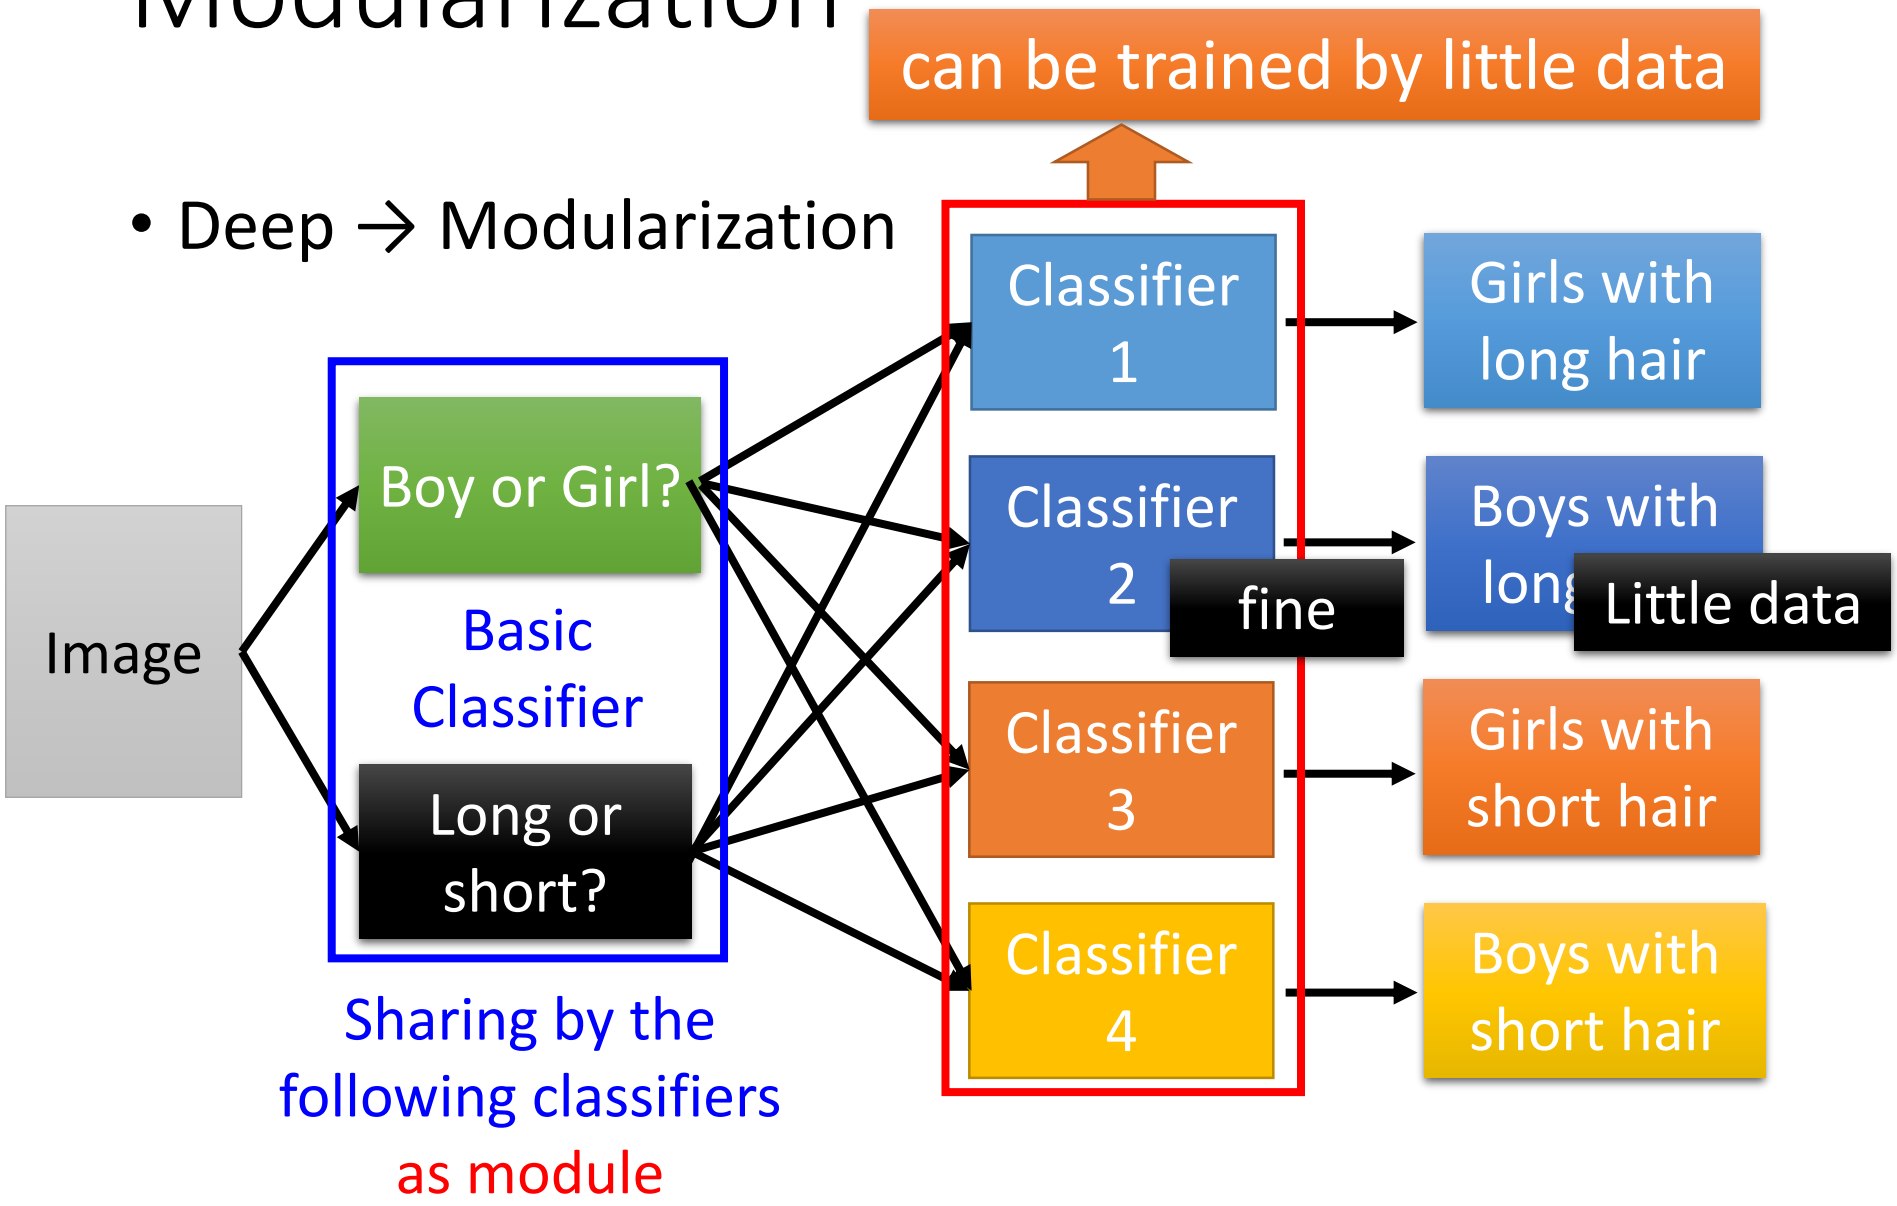

Modularization - Image

- Deep → Modularization

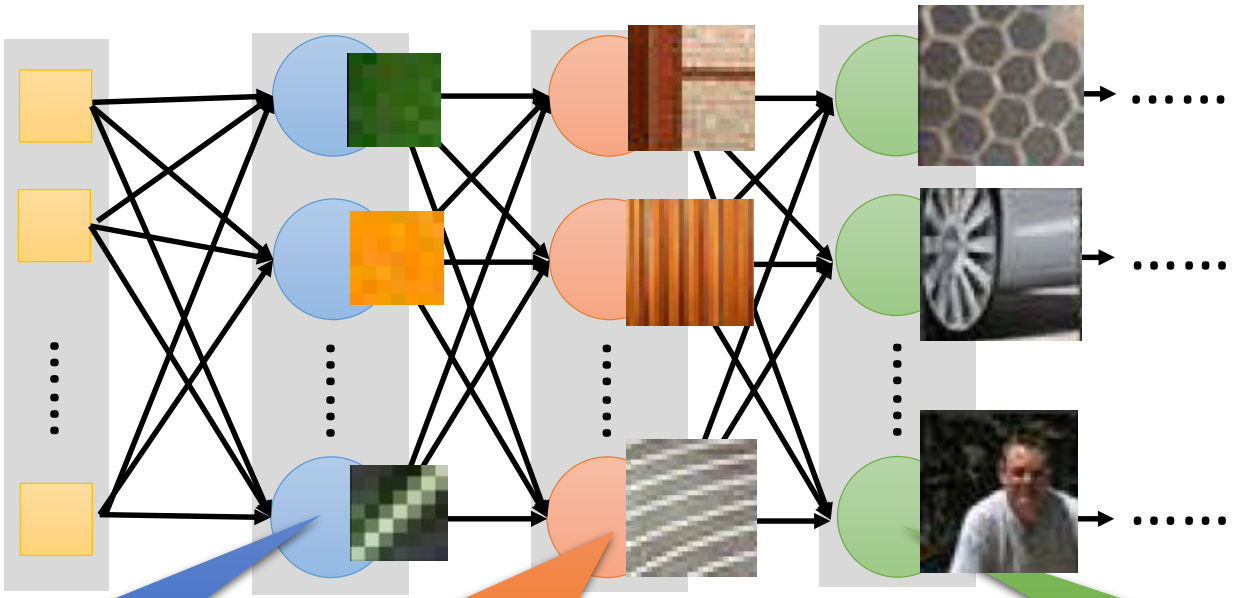

The most basic classifiers

Use 1st layer as module to build classifiers

Use 2nd layer as module ……

Reference: Zeiler, M. D., & Fergus, R. (2014). Visualizing and understanding convolutional networks. In *Computer Vision–ECCV 2014* (pp. 818-833)

Universality Theorem

Any continuous function f

$$f : R^N \rightarrow R^M$$

Can be realized by a network with one hidden layer

(given **enough** hidden neurons)

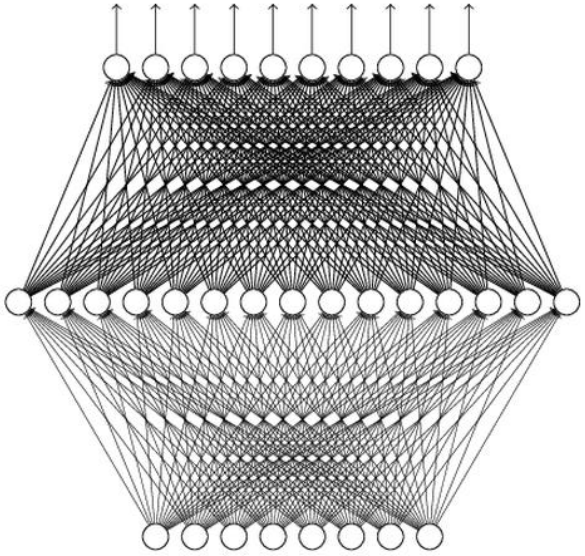

Reference for the reason:
http://neuralnetworksanddeeplearning.com/chap4.html

Yes, shallow network can represent any function.

However, using deep structure is more effective.

Analogy

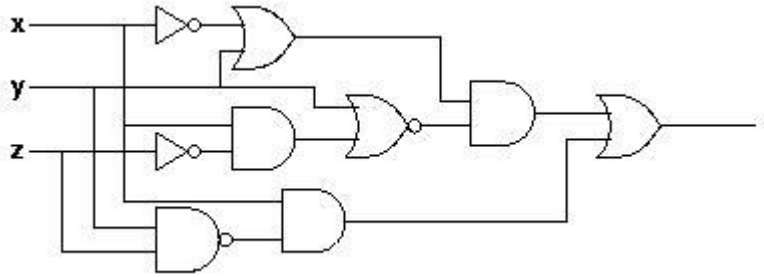

Logic circuits

- Logic circuits consists of **gates**
- **A two layers of logic gates** can represent **any Boolean function.**
- Using multiple layers of logic gates to build some functions are much simpler

→ less gates needed

Neural network

- Neural network consists of **neurons**
- **A hidden layer network** can represent **any continuous function.**
- Using multiple layers of neurons to represent some functions are much simpler

→ less parameters → less data?

This page is for EE background.

Analogy

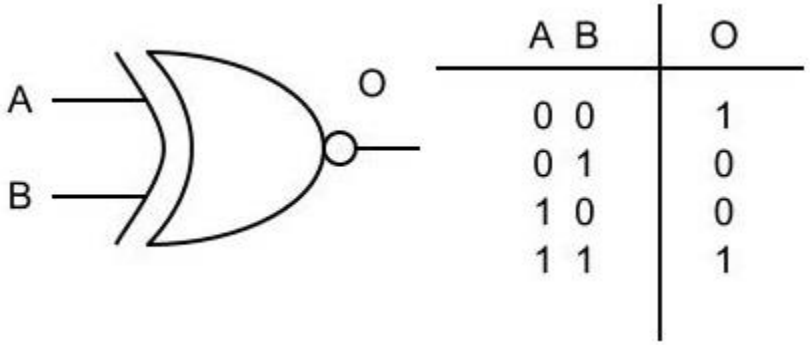

- E.g. ***parity check***

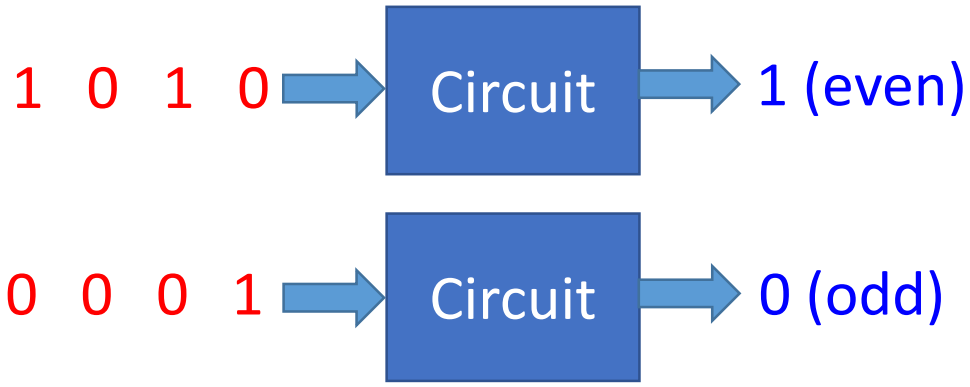

For input sequence with d bits,

Two-layer circuit need $O(2^d)$ gates.

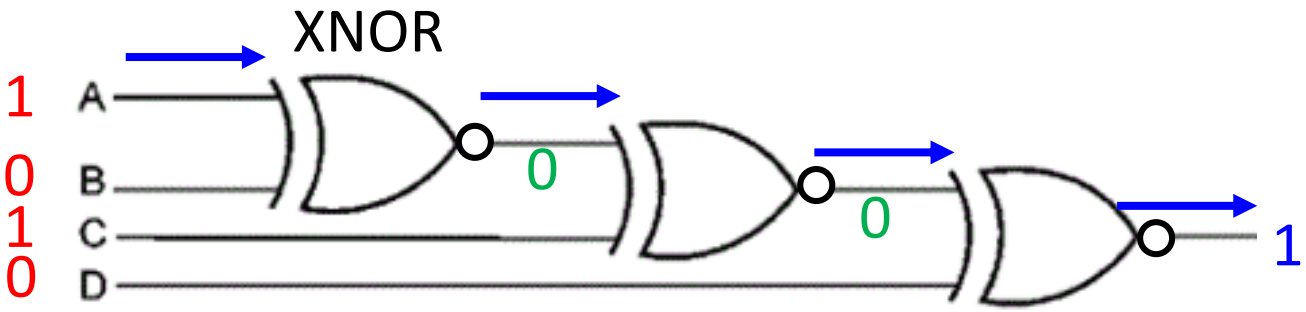

With multiple layers, we need only O(d) gates.

More Analogy

Use data effectively

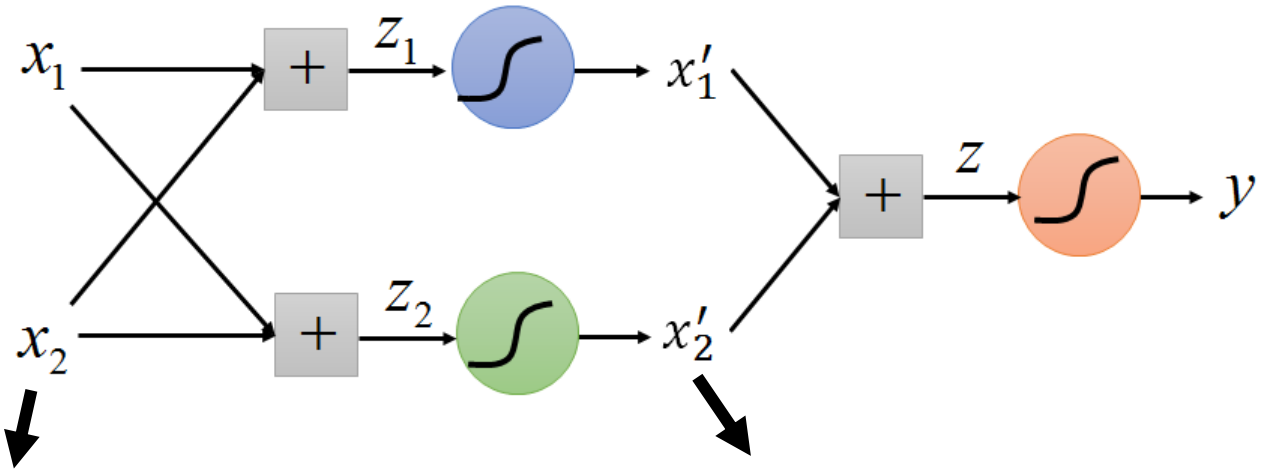

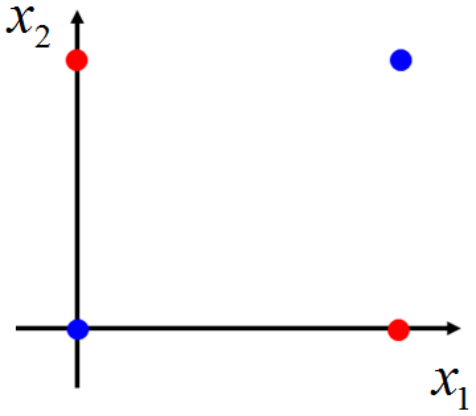

Folding the space

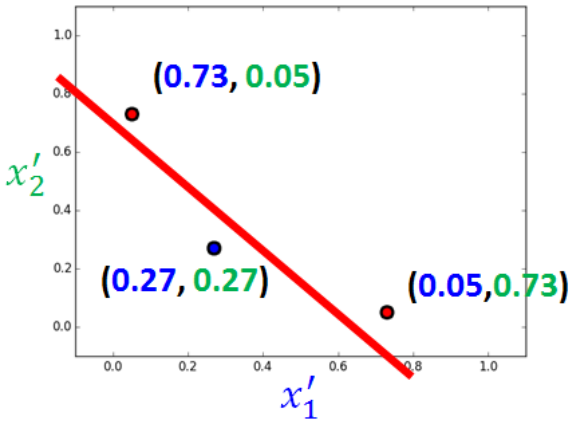

More Analogy - Experiment

Different numbers of training examples

10,0000 2,0000

$f : R^2 \rightarrow \{0,1\}$

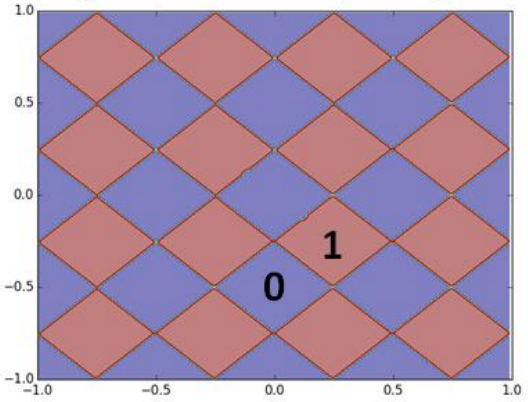

1 hidden layer

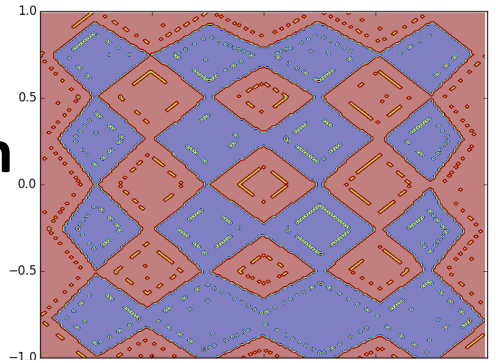

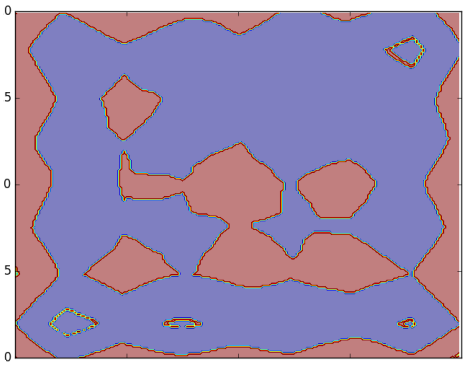

3 hidden layers

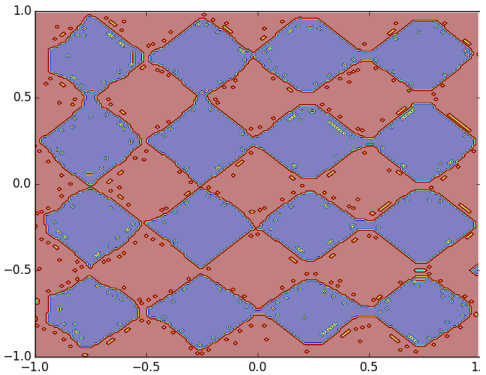

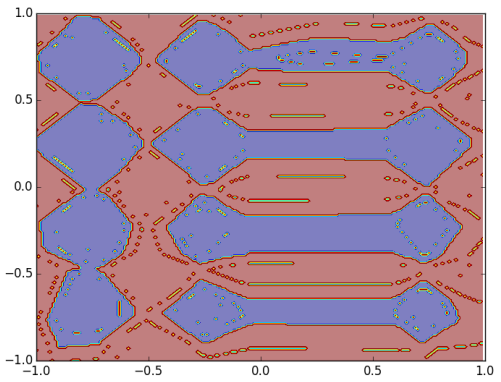

End-to-end Learning

- Production line

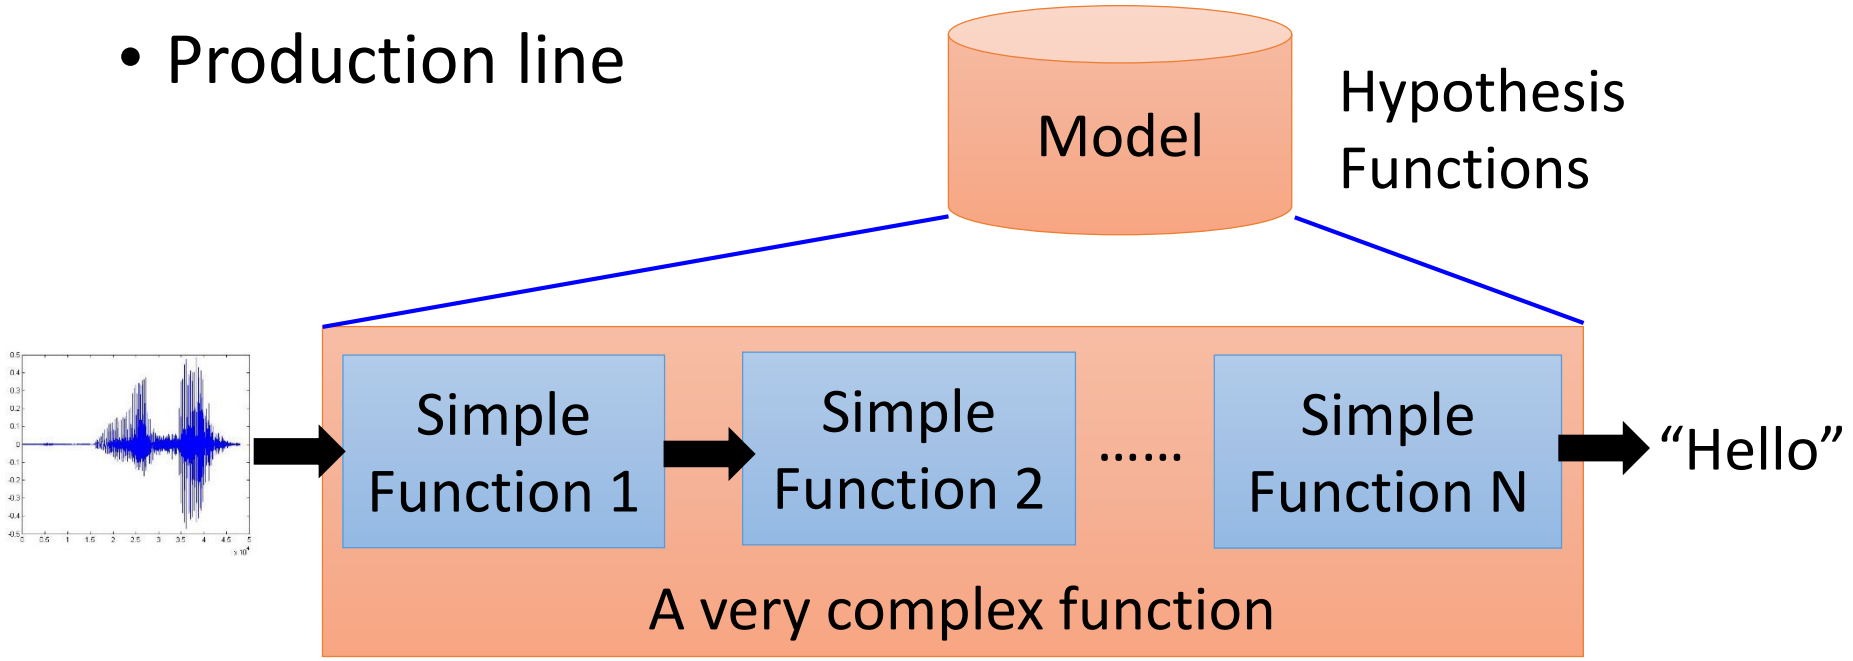

End-to-end training:

What each function should do is learned automatically

End-to-end Learning - Speech Recognition

- Deep Learning

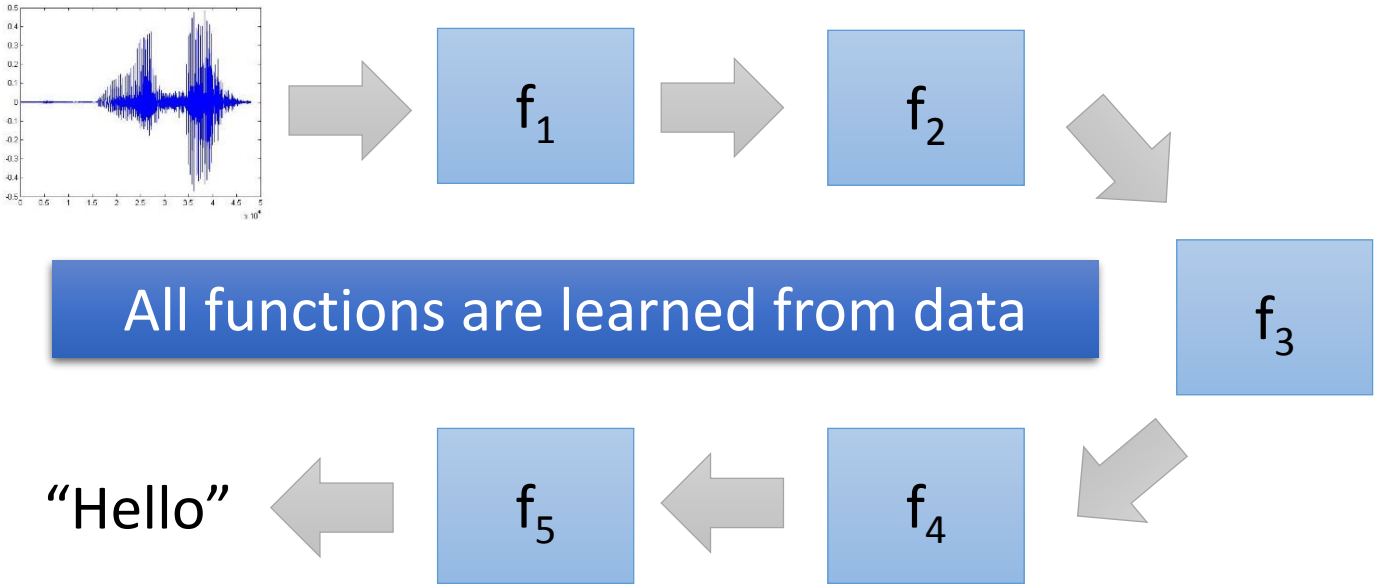

Less engineering labor, but machine learns more

Complex Task …

A. Mohamed, G. Hinton, and G. Penn, "Understanding how Deep Belief Networks Perform Acoustic Modelling," in ICASSP, 2012.

- Speech recognition: Speaker normalization is automatically done in DNN

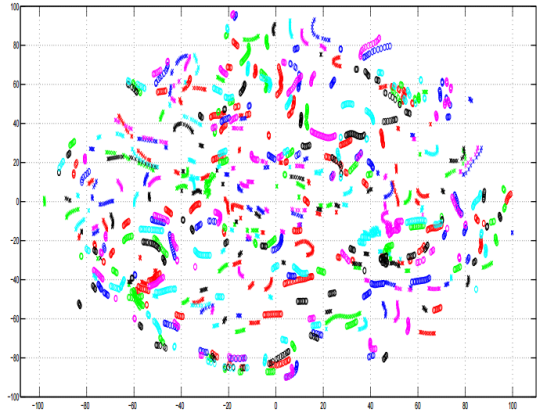

Input Acoustic Feature (MFCC)

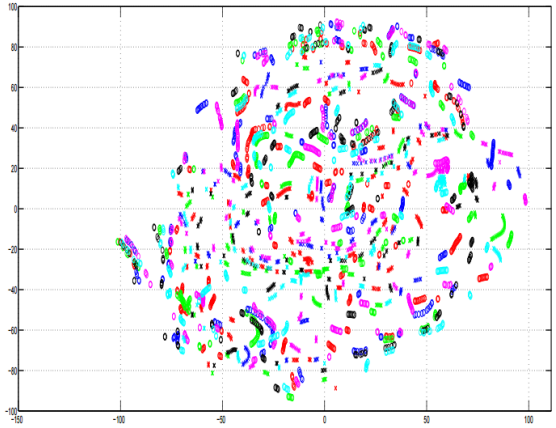

1-st Hidden Layer

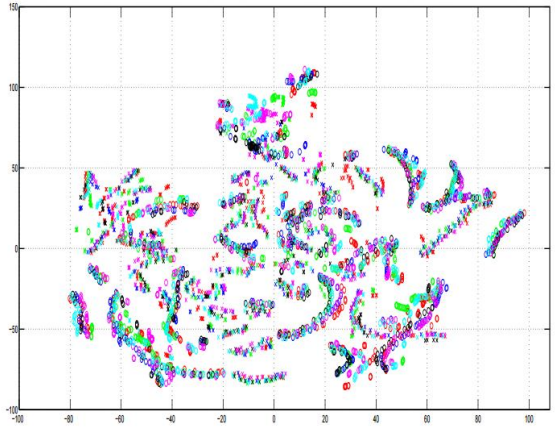

8-th Hidden Layer