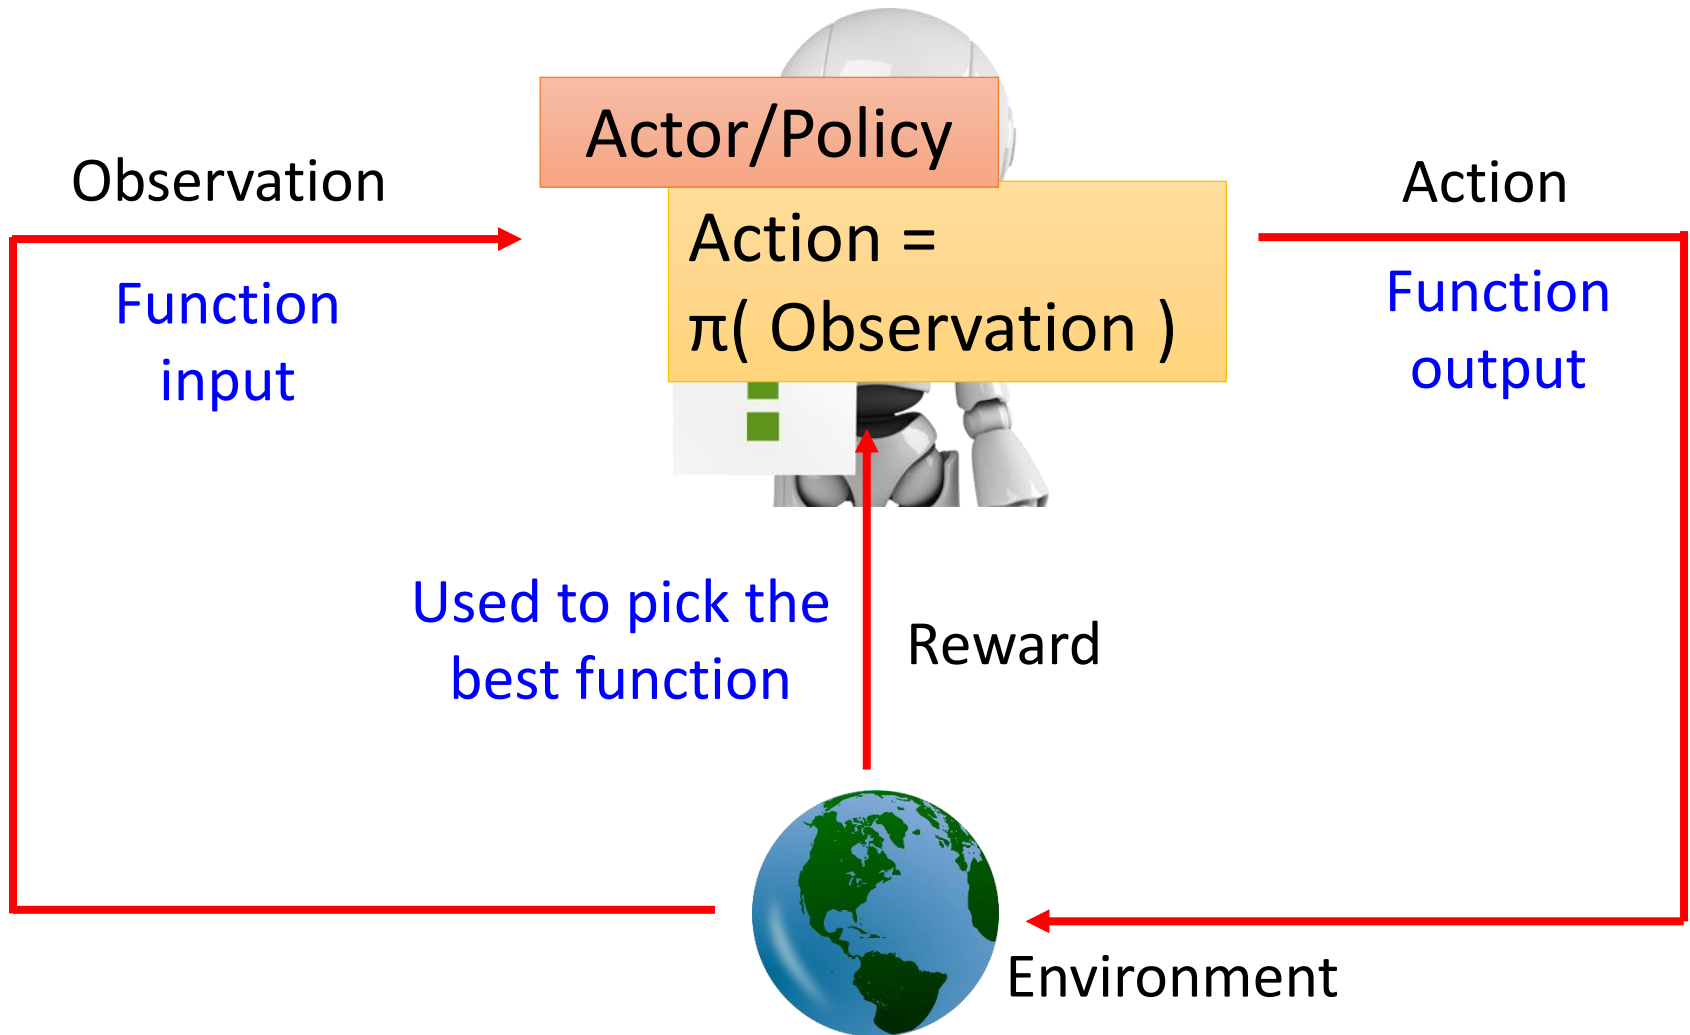


# Introduction of Reinforcement Learning

# Machine Learning $\approx$ Looking for a Func



# Properties of Reinforcement Learning

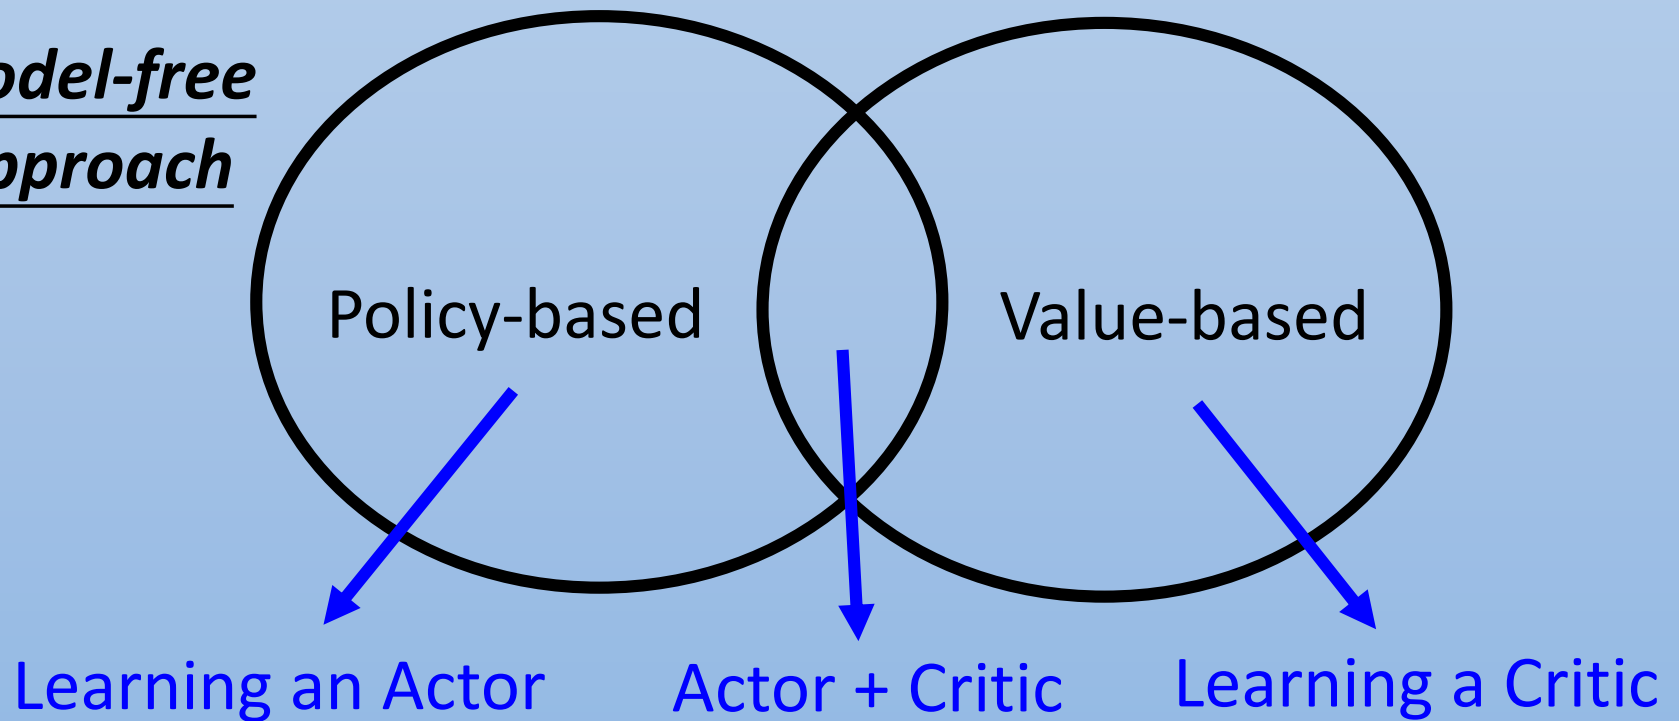
- **Reward delay**
  - In space invader, only “fire” obtains reward
    - Although the moving before “fire” is important
  - In Go playing, it may be better to sacrifice immediate reward to gain more long-term reward
- Agent's actions **affect the subsequent data it receives**
  - E.g. Exploration



# Outline

Alpha Go: policy-based + value-based  
+ model-based

## Model-free Approach



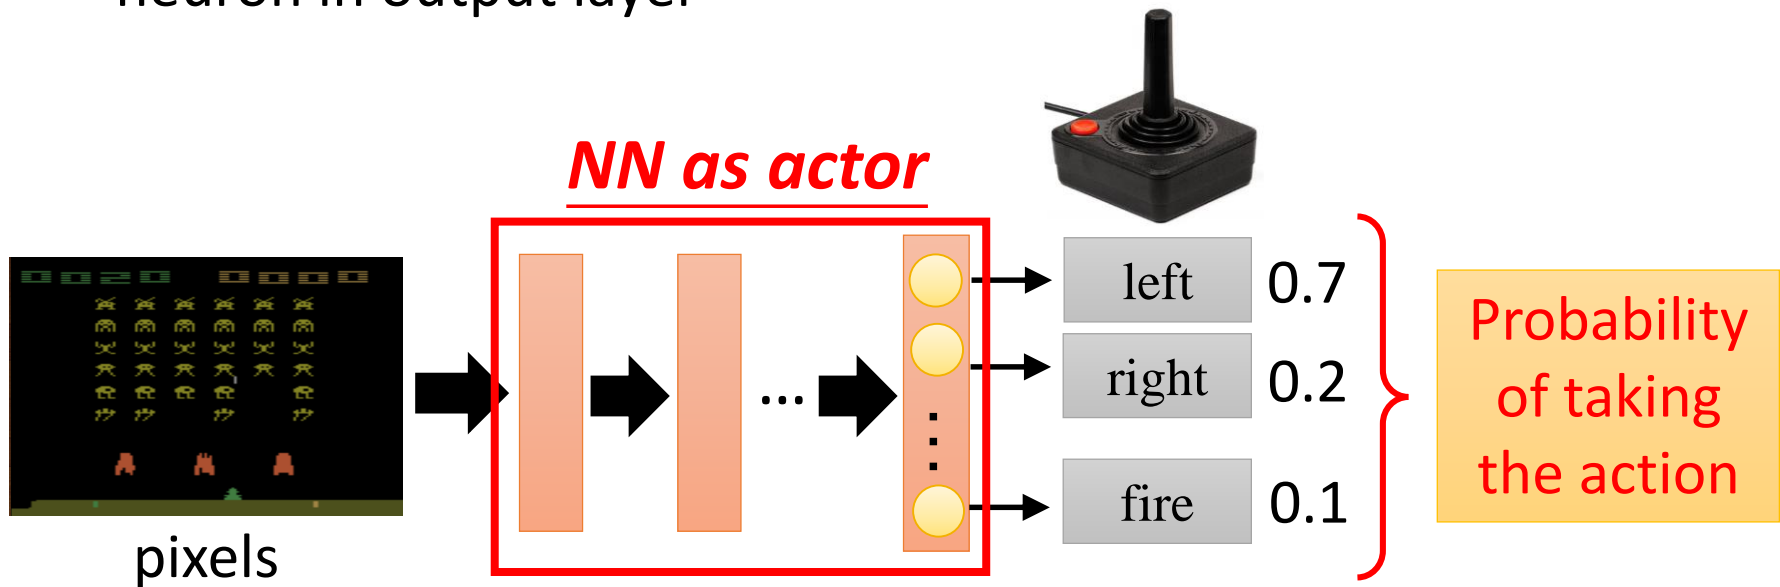
## Model-based Approach

# Policy-based Approach

## Learning an Actor

# Neural network as Actor

- Input of neural network: the observation of machine represented as a vector or a matrix
- Output neural network : each action corresponds to a neuron in output layer



What is the benefit of using network instead of lookup table?

generalization

# Goodness of Actor

- Given an actor  $\pi_{\theta}(s)$  with network parameter  $\theta$
- Use the actor  $\pi_{\theta}(s)$  to play the video game

We define  $\bar{R}_{\theta}$  as the expected value of  $R_{\theta}$

$\bar{R}_{\theta}$  evaluates the goodness of an actor  $\pi_{\theta}(s)$

- An episode is considered as a trajectory  $\tau$ 
  - $\tau = \{s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_T, r_T\}$
  - $R(\tau) = \sum_{t=1}^T r_t$
  - If you use an actor to play the game, each  $\tau$  has a probability to be sampled
    - The probability depends on actor parameter  $\theta$ :  
 $P(\tau|\theta)$

$$\bar{R}_{\theta} = \sum_{\tau} R(\tau)P(\tau|\theta) \approx \frac{1}{N} \sum_{n=1}^N R(\tau^n)$$

Sum over all  
possible trajectory

Total reward:  $R_{\theta} = \sum_{t=1}^T r_t$ ,

Even with the same actor

$R_{\theta}$  is different each time

Use  $\pi_{\theta}$  to play the

game  $N$  times,  
obtain  $\{\tau^1, \tau^2, \dots, \tau^N\}$

Sampling  $\tau$  from  $P(\tau|\theta)$   
 $N$  times

# Gradient Ascent

- Problem statement

$$\theta^* = \arg \max_{\theta} \bar{R}_{\theta}$$

- Gradient ascent

- Start with  $\theta^0$

- $\theta^1 \leftarrow \theta^0 + \eta \nabla \bar{R}_{\theta^0}$

- $\theta^2 \leftarrow \theta^1 + \eta \nabla \bar{R}_{\theta^1}$

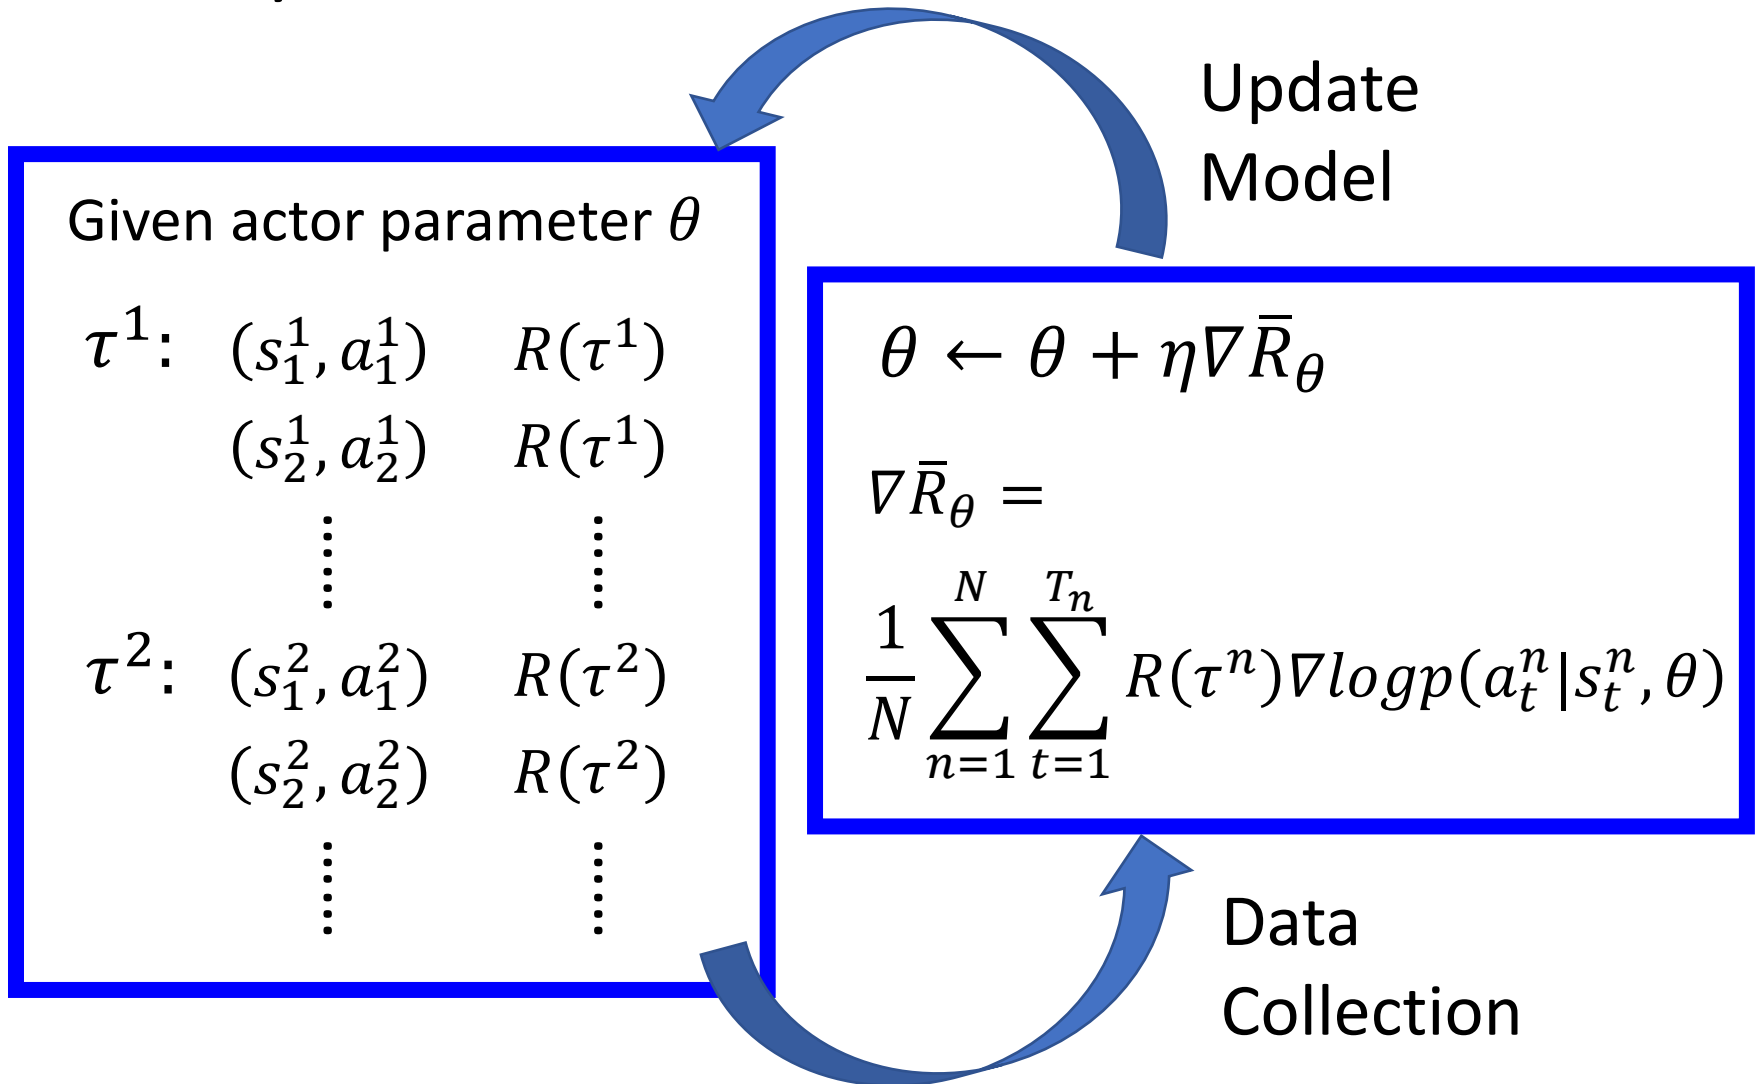
- .....

$$\theta = \{w_1, w_2, \dots, b_1, \dots\}$$

$$\nabla \bar{R}_{\theta} = \begin{bmatrix} \partial \bar{R}_{\theta} / \partial w_1 \\ \partial \bar{R}_{\theta} / \partial w_2 \\ \vdots \\ \partial \bar{R}_{\theta} / \partial b_1 \\ \vdots \end{bmatrix}$$



# Policy Gradient





# Add a Baseline

It is possible that  $R(\tau^n)$  is always positive.


$$\theta^{new} \leftarrow \theta^{old} + \eta \nabla \bar{R}_{\theta^{old}}$$

$$\nabla \bar{R}_{\theta} \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} (R(\tau^n) - \underline{b}) \nabla \log p(a_t^n | s_t^n, \theta)$$



Ideal  
case



It is probability ...



Sampling  
.....



The probability of the actions not sampled will decrease.




# Value-based Approach

## Learning a Critic

# Critic

- State value function  $V^\pi(s)$ 
  - When using actor  $\pi$ , the *cumulated* reward expects to be obtained after seeing observation (state)  $s$



$V^\pi(s)$  is large




$V^\pi(s)$  is smaller

# How to estimate $V^\pi(s)$

- Monte-Carlo based approach
  - The critic watches  $\pi$  playing the game


After seeing  $s_a$ ,

Until the end of the episode,  
the cumulated reward is  $G_a$




After seeing  $s_b$ ,

Until the end of the episode,  
the cumulated reward is  $G_b$



# How to estimate $V^\pi(s)$


- Temporal-difference approach



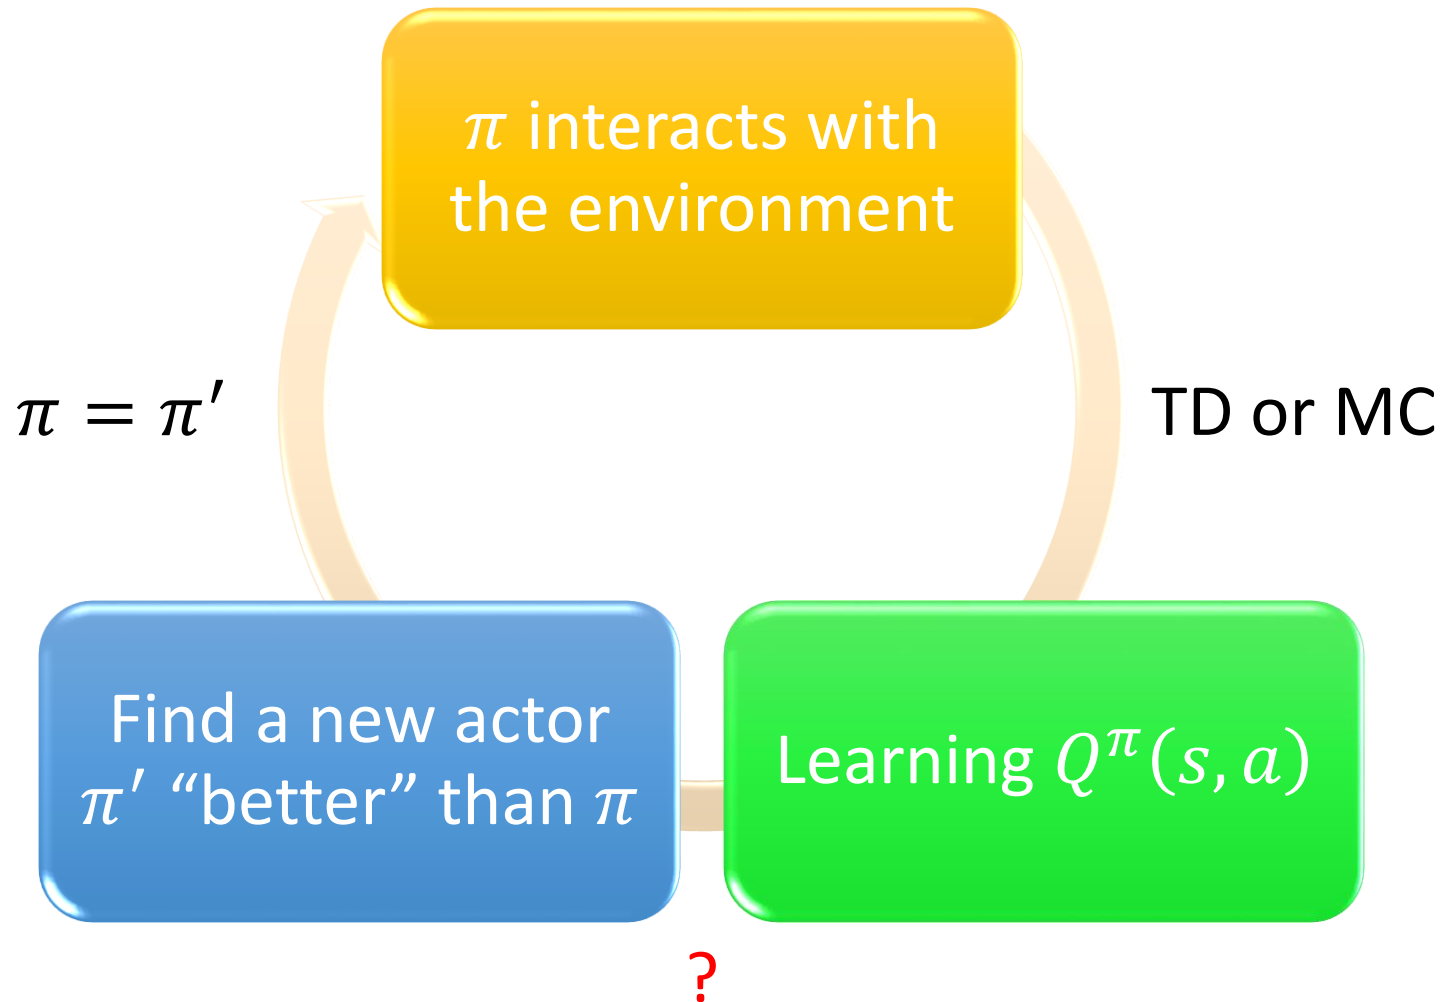
Some applications have very long episodes, so that delaying all learning until an episode's end is too slow.

# Another Critic

- State-action value function  $Q^\pi(s, a)$ 
  - When using actor  $\pi$ , the *cumulated* reward expects to be obtained after seeing observation  $s$  and taking  $a$




# Q-Learning





# Q-Learning



- Given  $Q^\pi(s, a)$ , find a new actor  $\pi'$  "better" than  $\pi$ 
  - "Better":  $V^{\pi'}(s) \geq V^\pi(s)$ , for all state  $s$

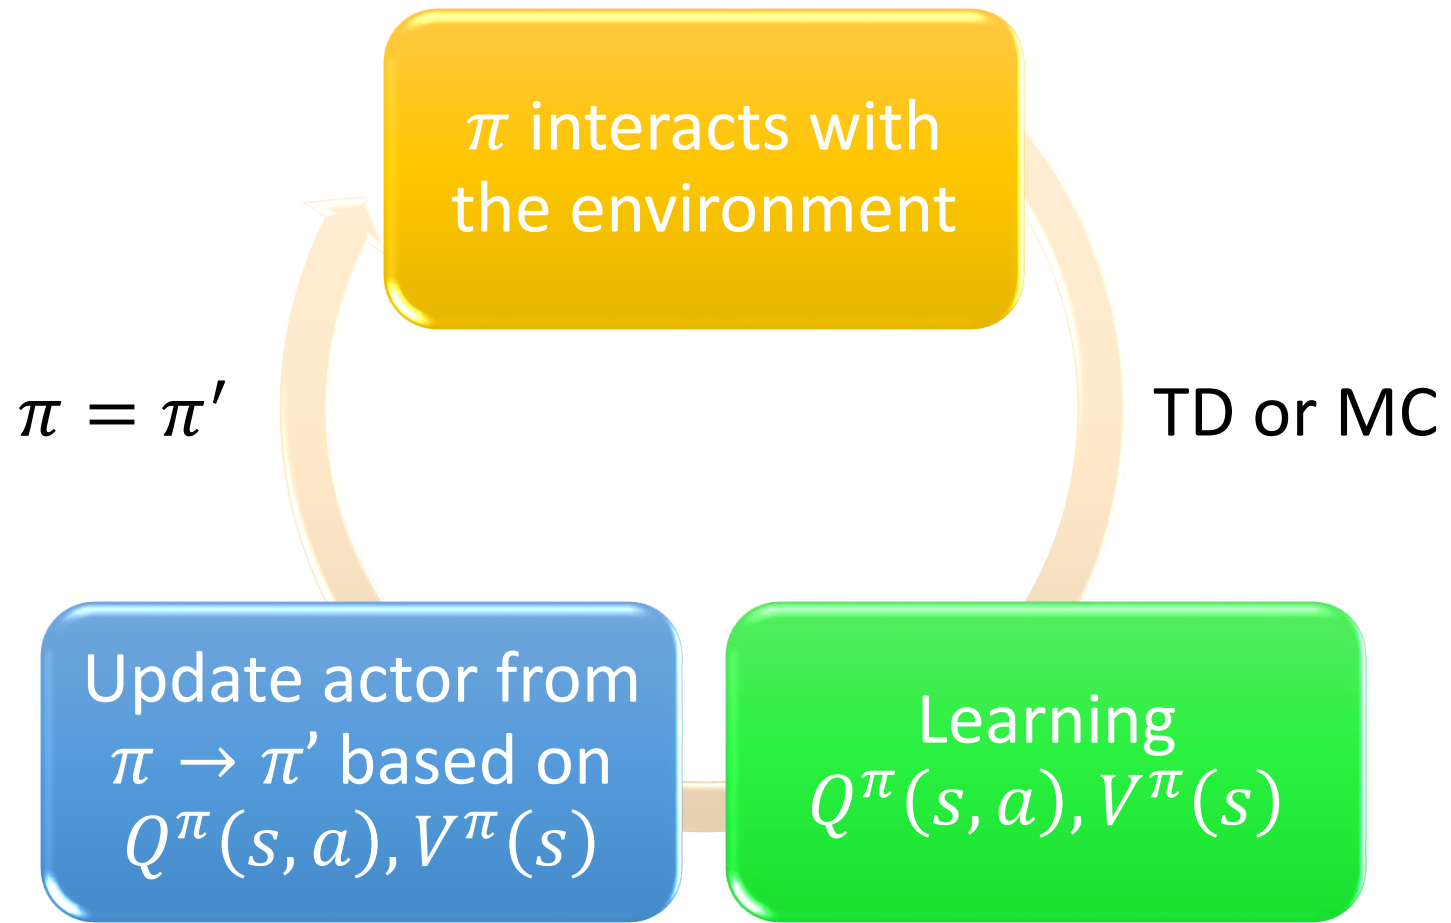
$$\pi'(s) = \arg \max_a Q^\pi(s, a)$$

- $\pi'$  does not have extra parameters. It depends on  $Q$
- Not suitable for continuous action  $a$

# Deep Reinforcement Learning


Actor-Critic

# Actor-Critic



# Actor-Critic

- Tips
  - The parameters of actor  $\pi(s)$  and critic  $V^\pi(s)$  can be shared



# Asynchronous

Source of image:

<https://medium.com/emergent-future/simple-reinforcement-learning-with-tensorflow-part-8-asynchronous-actor-critic-agents-a3c-c88f72a5e9f2#.68x6na7o9>

1. Copy global parameters
2. Sampling some data
3. Compute gradients
4. Update global models

